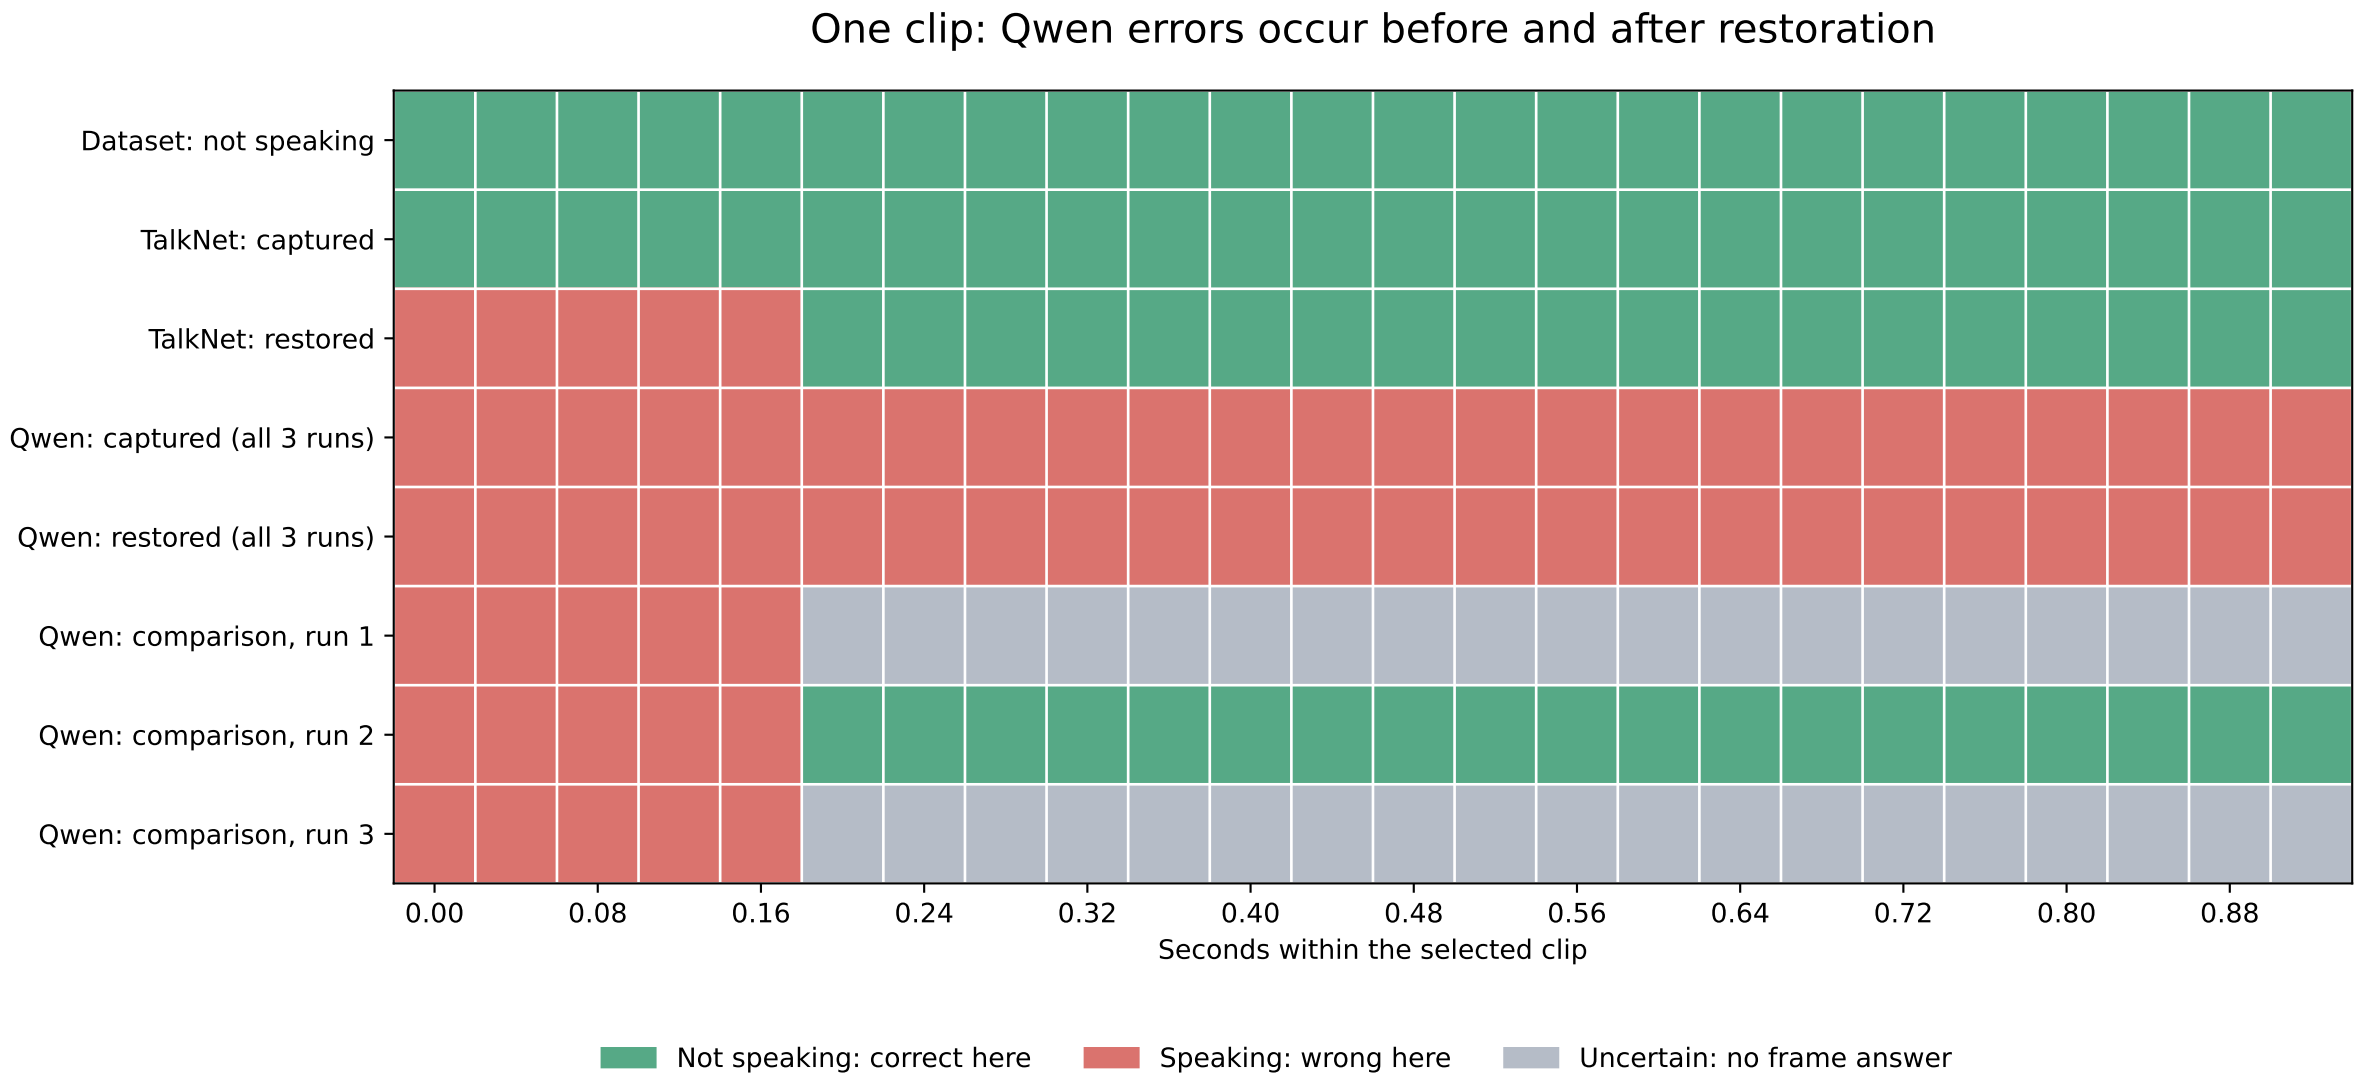
One clip: Qwen errors occur before and after restoration
Dataset: not speaking
TalkNet: captured
TalkNet: restored
Qwen: captured (all 3 runs)
Qwen: restored (all 3 runs)
Qwen: comparison, run 1
Qwen: comparison, run 2
Qwen: comparison, run 3
0.00
0.08
0.16
0.24
0.32
0.40
0.48
0.56
0.64
0.72
0.80
0.88
Seconds within the selected clip
Not speaking: correct here
Speaking: wrong here
Uncertain: no frame answer
Qwen uses tool measurements, not raw media. One targeted clip; repeated runs are not independent videos.