Qwen pilot: one wrong speaking answer; no frame-error improvement

8 completed tool-agent cases | 2 source videos | 4 tracks × 2 noise levels

1 / 8 wrong answers

35 / 192 selected frame decisions wrong

0.50 confidence on the wrong answer

A. Qwen's final clip answers

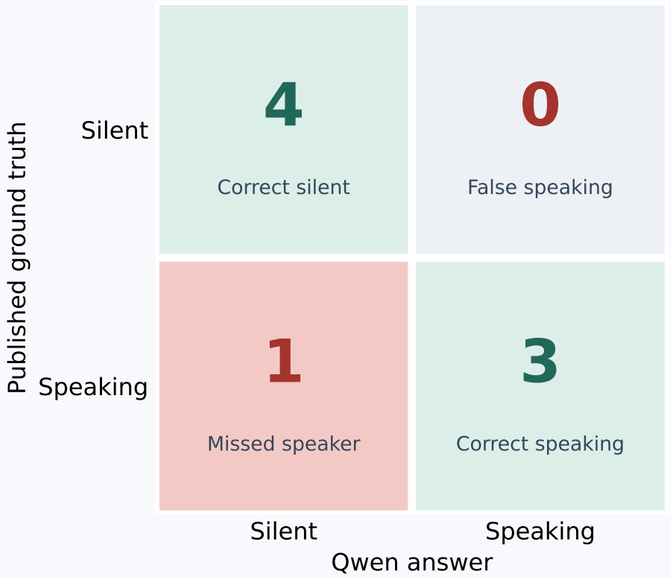

B. Errors in Qwen-selected evidence

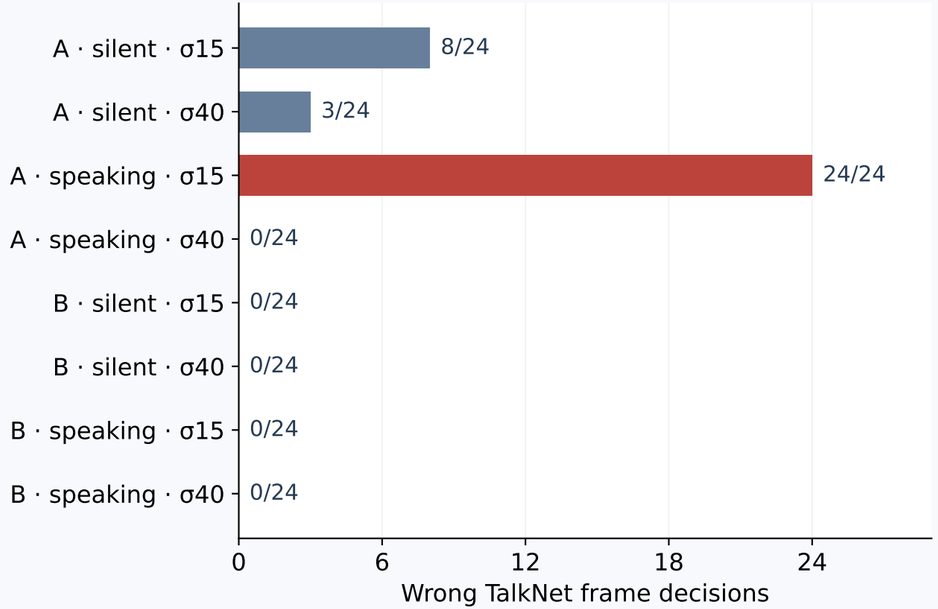

C. Selection baselines

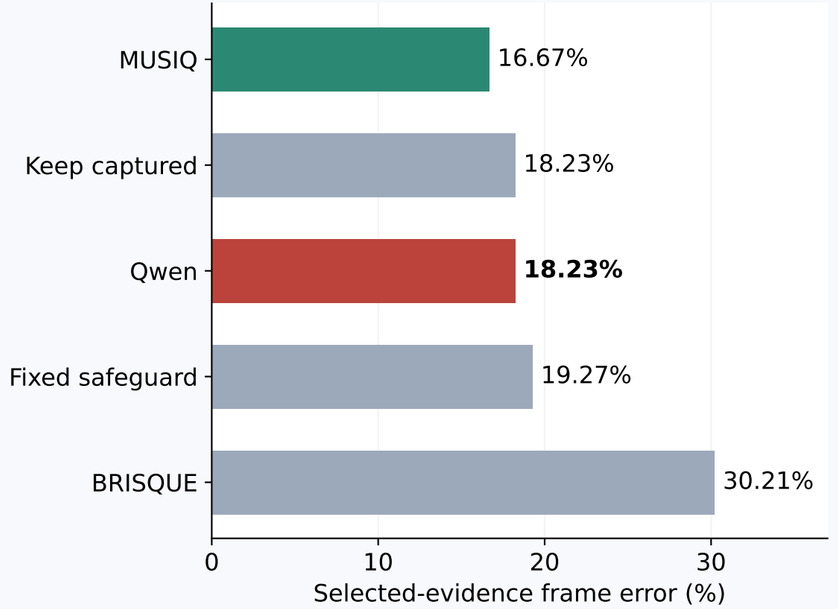

What failed: Qwen missed one speaking clip; the underlying detector also missed it in the clean reference.

What this does not show: a highly confident Qwen error, or proof that restoration caused its wrong clip answer.

Qwen used tool measurements, not direct video/audio. Panels B–C measure TalkNet errors in selected video, not Qwen frame predictions.

Small descriptive pilot: noise variants and adjacent frames are correlated. No population-level confidence interval is claimed.

Source: saved Qwen/Groq traces and CPU pilot measurements • 15 October 2026